

A Verified Garbage Collector for Gallina

Shengyi Wang[†], Anshuman Mohan[†], Qinxiang Cao[‡], Aquinas Hobor[†]



APLAS NIER
December 1, 2019

Verify **graph-manipulating** programs
written in **executable C**
with **machine-checked** correctness proofs

Hard, but ubiquitous in critical areas



Certifying Graph-Manipulating C Programs via Localizations within Data Structures

SHENGYI WANG, National University of Singapore, Singapore

QINXIANG CAO, Shanghai Jiao Tong University, China

ANSHUMAN MOHAN, National University of Singapore, Singapore

AQUINAS HOBOR, National University of Singapore, Singapore

VST + CompCert + 25000 LOC library

Powerful enough to verify **executable code**
against **realistic specifications**
expressed with **mathematical graphs**

[Wang *et. al.*, PACMPL OOPSLA 2019]



Gallina \rightsquigarrow CompCert C \rightsquigarrow Assembly

Gallina assumes **infinite** memory
but CompCert C has a **finite** heap


Solution: garbage collect the CompCert C code

New problem: verify the garbage collector

Our Garbage Collector

GC has jurisdiction over the heap

Mutator **allocs** in special subheap




Our Garbage Collector

GC has jurisdiction over the heap

Mutator **allocs** in special subheap

If subheap is full




Our Garbage Collector

GC has jurisdiction over the heap

Mutator **allocs** in special subheap


If subheap is full **call GC** and try again



- 12 generations, doubling in size
- Functional mutator: no back pointers
- Cheney's mark-and-copy collects gen to next
- Potentially triggers cascade of pairwise collections
- Three key functions:
 - `forward` copies individual objects
 - `do_scan` repairs copied objects
 - `forward_roots` kick-starts the collection


Intuitive Specification

Primum non nocere: first, do no harm




Overview of Operations

Nursery cannot fit `alloc`



Overview of Operations

Nursery cannot fit `alloc`
`do_gen`




Overview of Operations

Nursery cannot fit alloc

do_gen

forward_roots




Overview of Operations

Nursery cannot fit alloc

do_gen

forward_roots


forward



Overview of Operations

Nursery cannot fit alloc


```
do_gen  
  forward_roots  
    forward  
do_scan
```



Overview of Operations

Nursery cannot fit alloc


```
do_gen  
  forward_roots  
    forward  
do_scan  
  forward
```



Overview of Operations

Nursery cannot fit alloc


```
do_gen
  forward_roots
    forward
do_scan
  forward
reset_gen
```



Overview of Operations


Nursery cannot fit alloc

```
do_gen  
  forward_roots  
    forward  
do_scan  
  forward  
reset_gen
```



Overview of Operations


Non-Concerns
more garbage



Overview of Operations

Non-Concerns

- more garbage
- backward pointers




Overview of Operations

Non-Concerns

- more garbage
- backward pointers


Sources of Complexity

- variable-length objects
- disambiguate int/ptr
- determine v 's gen
- determine gen size
- what if `malloc` fails?



Instantiating GC_Graph

A PreGraph is a hextuple ($VType$, $EType$, $vvalid$, $evalid$, src , dst)



A PreGraph is a hextuple (VType, EType, vvalid, evalid, src, dst)

GC_PreGraph: VType := nat * nat
EType := VType * nat
src := fst
dst := *unrestricted*
 $\forall v. \text{vvalid}(\gamma, v) \Leftrightarrow \text{graph_has_v}(\gamma, v)$
 $\forall v, out. \text{evalid}(\gamma, (v, out)) \Leftrightarrow$
 $\text{vvalid}(\gamma, v) \wedge \text{In } out \text{ (get_edges}(\gamma, v))$

Instantiating GC_Graph

A LabeledGraph is a quadruple (PreGraph, VL, EL, GL)

GC_Graph: GC_PreGraph as shown

VL := raw_vert_block

EL := unit

GL := list gen_info

Definition

raw_fld := Z + GC_Ptr.

Record raw_vert_block :=
{ raw_mark: bool;
 copied_vertex: VType;
 raw_flds: list raw_fld;
 (* elided *) }.

Record gen_info :=
{ s_addr: val;
 s_ok: isptr s_addr;
 num_vert: nat;
 (* elided *) }.

forward: a Deep Dive

forward is everywhere!

forward is **robust**

pointer? in from space? already forwarded?
and **versatile**
called on root set called on heap

```
void forward (value *s, *l, **n, *p) {  
  value * v; value va = *p;  
  if(Is_block(va)) {  
    v = (value*)iop2ptr(va);  
    if(Is_from(s, l, v)) {  
      header_t hd = Hd_val(v);  
      if(hd == 0) {  
        *p = Field(v,0);  
      } else { /* elided */  

```


$$\left\{ \begin{array}{l} \forall \gamma, from, to, v, n. \text{gc_graph}(\gamma) \wedge \text{compat}(\gamma, from, to) \wedge \\ s = \text{start}(\gamma, from) \wedge l = s + \text{gensz}(\gamma, from) \wedge \\ n = \text{nxtaddr}(to) \wedge p = \text{vaddr}(\gamma, v) + n \end{array} \right\} \stackrel{\text{def}}{=} \phi_1$$

```
void forward (value *s, *l, **n, *p) {
```

```
  /* elided */
```

```
  if(hd == 0) {
```

```
    *p = Field(v,0);
```

$$\left\{ \begin{array}{l} \phi_1 \wedge \exists \gamma'. \text{gc_graph}(\gamma') \wedge \gamma' = \text{upd_edge}(\gamma, e, \text{copy}(\gamma, v)) \wedge \\ \text{compat}(\gamma', from, to) \wedge \text{fwd_relation}(\gamma, \gamma', from, to, v, n) \end{array} \right\}$$

```
}
```

```

else {
  int i; int sz; value *new; sz = size(hd);
  new = *next+1; *next = new+sz; Hd_val(new) = hd;
  for(i = 0; i < sz; i++)
    Field(new, i) = Field(v, i);
  {
 $\phi_1 \wedge \exists \gamma', v'. \text{gc\_graph}(\gamma') \wedge v' = \text{copied\_vertex}(\gamma, to) \wedge$ 
 $\gamma' = \text{copy\_vertex}(\gamma, to, v, v') \wedge \text{compat}(\gamma', from, to)$ 
 $\} \stackrel{\text{def}}{=} \phi_2$ 
  Hd_val(v) = 0; Field(v, 0) = p2iop((void *)new);
  *p = p2iop((void *)new);
  {
 $\phi_2 \wedge \exists \gamma''. \text{gc\_graph}(\gamma'') \wedge \gamma'' = \text{upd\_edge}(\gamma', e, v') \wedge$ 
 $\text{compat}(\gamma'', from, to) \wedge \text{fwd\_relation}(\gamma, \gamma'', from, to, v, n)$ 
 $\}$ 
}

```

```
Inductive fwd_relation from to :  
  forward_t -> LGraph -> LGraph -> Prop :=  
| fr_v_not_in : forall v g,  
  vgen v <> from ->  
  fwd_relation from to (inl (inr v)) g g  
| fr_e_to_fwded : forall e g,  
  vgen (dst g e) = from ->  
  raw_mark (vlabel g (dst g e)) = true ->  
  let new_g := lgraph_gen_dst g e  
    (copied_vertex (vlabel g (dst g e))) in  
  fwd_relation from to (inr e) g new_g
```

```
| fr_e_to_not_fwded_Sn : forall e g g',  
  vgen (dst g e) = from ->  
  raw_mark (vlabel g (dst g e)) = false ->  
  let new_g :=  
    lgraph_gen_dst (lgraph_copy1v g (dst g e) to)  
      e (copy1v_new_v g to) in  
  fwd_loop from to  
    (make_fields new_g (copy1v_new_v g to)) new_g g' ->  
  fwd_relation from to (inr e) g g'
```

Similar to `forward_relation`, we have

`forward_roots_relation`

`do_scan_relation`

`do_generation_relation`


`garbage_collect_relation`

A composition of these gives us our **isomorphism**

Moving Towards Isomorphism

But the journey is far from easy!

A brief look at `semi_iso`:



The general iterative strategy:

$$\frac{\frac{}{\gamma \triangleright \gamma} \quad \frac{\gamma \triangleright \gamma_i \quad \gamma_i \rightsquigarrow \gamma_{i+1}}{\gamma \triangleright \gamma_{i+1}}}{\gamma_\alpha \triangleright \gamma_\omega}$$

A specific example:

```
Lemma semi_iso_refl: forall g from to,  
  sound_gc_graph g -> semi_iso g g from to nil.
```

```
Lemma fwd_rel_semi_iso:  
  forall from to p g1 g2 g3 roots,  
    semi_iso g1 g2 from to l1 ->  
    forward_relation from to p g2 g3 ->  
    semi_iso g1 g3 from to
```


And eventually,

```
Theorem garbage_collect_iso: forall roots1 roots2 g1 g2,  
  ...  
  garbage_collect_relation roots1 roots2 g1 g2 ->  
  gc_graph_iso g1 roots1 g2 roots2.
```

The graphs are isomorphic

up to the vertices reachable from roots


The space between `n` and `l` is available for `alloc`

Note that we may still not achieve full isomorphism:

the **graph label** may change to register new vertices
and may even grow to accommodate new generations

Recap: Intuitive Specification

Primum non nocere: first, do no harm



- Cheney implemented too conservatively:
only part of `to` space needs to be
scanned during `do_scan`

Performance doubled


- Overflow in the following calculation:

```
int space_size =  
    h->spaces[i].limit - h->spaces[i].start;  
if difference > 231
```

Fixed by adjusting nursery size

Double-bounded pointer comparisons:

```
int Is_from(value * lo, value * hi, value * v) {  
    return (lo <= v && v < hi);  
}
```



Resolved using CompCert's `extcall_properties`

A classic OCaml trick to disambiguate int/ptr:

```
int test_int_or_ptr (value x) {  
    return (int)((((intnat)x)&1); }
```

Essentially, assume that pointers are even-aligned.

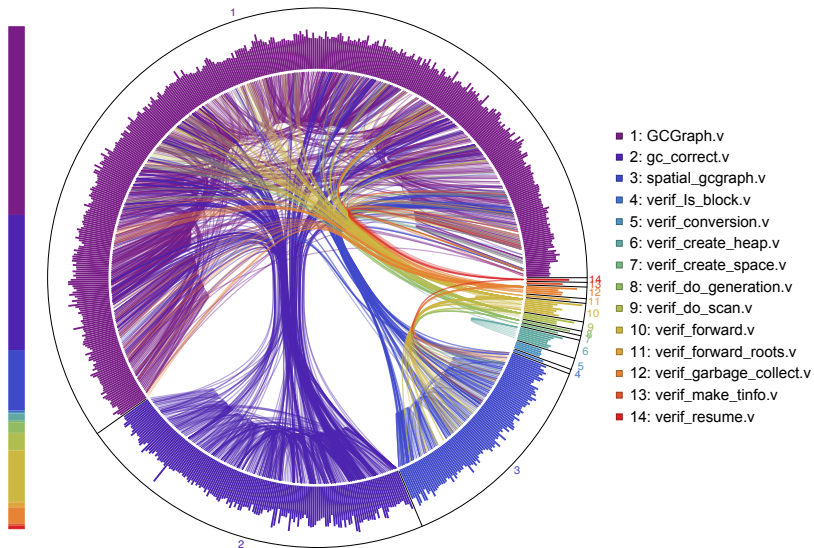
Consider:

```
void foo() {  
    char a; char b; char* pa = &a; char* pb = &b;  
    if ((pa&1 == 0) && (pb&1 == 0)) { /* elided */ } }
```

True in C, false in exec!

Discussing char alignment issues with CompCert

Reusability: separation between pure and spatial reasoning



Problems of a **similar shape**
serialization
other collectors

Towards a verified GC for **OCaml**
mutability
calculate root set
allow other datatypes

Further refinements required in C semantics
before we can **specify** and **verify** OCaml's GC?